


# Small String Optimization

## What are strings/why optimize?

- strings: of characters!
  - websites you browse
  - name and passwords for bank transactions
  - character played in game
  - etc etc...
- modern databases store LOTS of data; we need the performance!

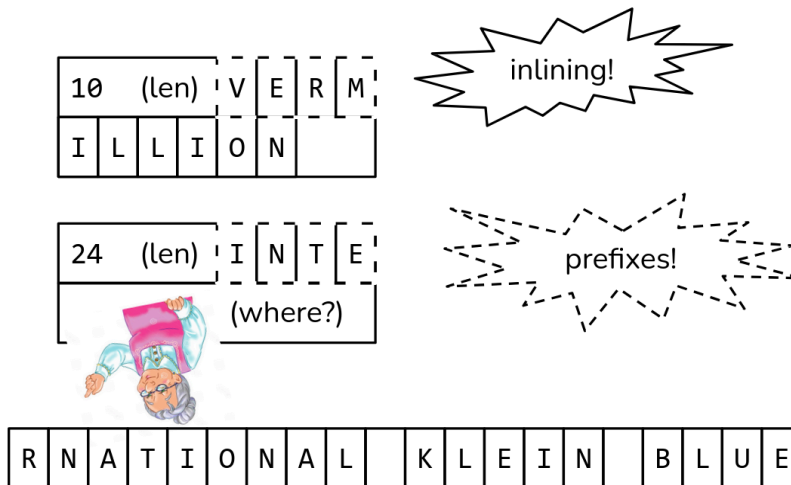


## Technical overview of strings



- following the pointer grandma is slow!

## How we can optimize



- inlining: in short strings, put the string where the pointer is. we never use the pointer grandma!
- prefixes: operations like string comparison don't need to follow pointer grandma until the 5th letter

## Further Reading

[https://cedardb.com/blog/german\\_strings/](https://cedardb.com/blog/german_strings/)

<https://pola.rs/posts/polars-string-type/>

## Asides

What is a byte?  
Computers store information in 1s and 0s. Each 1 or 0 is a “bit.” There are eight bits in a byte!

Four bits are in a nibble :)

Why did we choose a structure 16 bytes large to represent a string? It's just a magic upper bound; larger string structures can be slower. See [https://cedardb.com/blog/strings\\_deep\\_dive/#function-call-optimizations](https://cedardb.com/blog/strings_deep_dive/#function-call-optimizations).